

Rishimithan Kannan

Chennai, Tamil Nadu, India | rishimithan@gmail.com | +91 7010734901 | github.com/Rishimithan10
linkedin.com/in/rishimithan-kannan

Professional Summary

Data Analyst and Data Scientist specializing in Azure cloud data pipelines and ML-driven analytics, with a strong foundation in statistical modeling and business intelligence. Experienced in building end-to-end Lakehouse architectures, forecasting models, and Power BI dashboards that translate complex data into actionable insights.

Technical Skills

Languages: Python, SQL, C++, CUDA

Cloud & Data Platforms: Azure (ADF, ADLS Gen2, Databricks, Synapse Analytics), AWS (Lambda, SNS, SES)

Analytics & ML: Scikit-learn, Pandas, NumPy, ARIMA/SARIMA, Prophet, XGBoost, PySpark, PyTorch, HuggingFace Transformers

BI & Visualization: Power BI, IBM Cognos, Matplotlib, Seaborn

Databases: MySQL, Azure SQL, Pinecone Vector DB, Delta Lake

Tools: Git, Jupyter, SHAP

Work Experience

AI Research Intern · BluBridge Technologies, Chennai Jan 2026 – May 2026

- Designed structured **benchmarking frameworks** to collect and analyze performance metrics (throughput, latency, memory efficiency) across multiple hardware configurations; translated profiling data into comparative analytical reports to support optimization decisions.
- Conducted **systematic performance analysis** across varying matrix workload sizes, tracking compute efficiency trends and quantifying improvements using structured measurement methodologies.

Graduate Engineering Trainee · PSIOG Digital (P) Ltd, Chennai Aug 2024 – Feb 2025

- Built and maintained **ETL pipelines** on **Azure SQL Database** to transform raw transactional data into structured analytical models with automated data validation, ensuring high data quality for reporting.
- Designed **Power BI dashboards** surfacing demand forecasts, KPIs, and operational metrics, enabling data-driven decisions for non-technical stakeholders across logistics business units.
- Trained **ARIMA and SARIMA time-series models** on **Azure ML** for vehicle demand and pricing prediction; integrated forecast outputs into analytical pipelines to support operational reporting.
- Architected **event-driven data workflows** using **AWS Lambda, SNS, and SES** for automated real-time data ingestion, reducing manual intervention and improving pipeline reliability.

Projects

Rossmann Store Sales – Demand Forecasting

<https://github.com/Rishimithan10/Rossmann-Store-Sales-Demand-Forecasting>

- Built an end-to-end demand forecasting pipeline on **1,017,209 sales records** across **1,115 stores**, comparing SARIMA, Prophet, and XGBoost on a 6-week holdout.
- Engineered **31 features** including lag variables (lag_1, lag_7, lag_28), rolling statistics, and calendar features; applied time-based train/test split to prevent data leakage.
- Trained a single global **XGBoost** model across all 1,115 stores achieving **5.91% MAPE** on Store 1 – outperforming SARIMA (8.49%) by 30% and Prophet (7.96%) by 26%.
- Achieved **under 10% MAPE on 776/1,115 stores** (69.6%) with a single global model; applied **SHAP** analysis to identify lag features and Promo as top sales drivers.

Fraud Analytics – Azure Lakehouse Pipeline

<https://github.com/Rishimithan10/azure-fraud-detection-pipeline>

- Architected a **3-layer Medallion Lakehouse** (Bronze → Silver → Gold) on **Azure Data Lake Storage Gen2**, establishing a scalable and structured data foundation for fraud analytics.

- Orchestrated end-to-end data movement using **Azure Data Factory (ADF)** with automated ingestion scheduling, dependency management, and error handling across pipeline stages.
- Engineered data quality layer in **Databricks (PySpark)** – performing null handling, type casting, deduplication, and feature normalization – ensuring clean, analysis-ready data in the Silver layer.
- Computed fraud-related **KPI metrics** in the Gold layer – fraud rate by merchant category, transaction volume trends, and anomaly scores – using **Delta tables** for versioned, reliable storage.
- Enabled self-serve fraud reporting across customer segments via **Azure Synapse Analytics (SQL Serverless)** analytical views, reducing ad-hoc query turnaround for business stakeholders.

Education

B.Tech in Electronics and Communication Engineering

2020 – 2024

SASTRA Deemed University, Thanjavur, Tamil Nadu – **CGPA: 7.92 / 10**

Certifications

- Deloitte Data Analytics Job Simulation Certificate (2025)
- Machine Learning Specialization – Andrew Ng, Coursera (2022)
- Python for Data Science, AI, and Development – Coursera (2022)
- Data Visualization with Excel and Cognos – Coursera (2022)